

```
In [153]: import pandas as pd
import seaborn as sns

from pandas.core.dtypes.common import is_numeric_dtype
import matplotlib.pyplot as plt
```

```
In [154]: #import dataset from github
bank_data = 'https://raw.githubusercontent.com/rashakil-ds/Public-Datasets/main/bank_data.csv'
```

```
In [155]: df = pd.read_csv(bank_data)
```


```
In [156]: #showing first 5 rows
df.head()
```

```
Out[156]:
```

	ID	Customer_ID	Month	Name	Age	SSN	Occupation	Annual_Income	Monthly
--	----	-------------	-------	------	-----	-----	------------	---------------	---------

0	0x160a	CUS_0xd40	September	Aaron Maashoh	23	821- 00- 0265	Scientist	19114.12	
1	0x160b	CUS_0xd40	October	Aaron Maashoh	24	821- 00- 0265	Scientist	19114.12	
2	0x160c	CUS_0xd40	November	Aaron Maashoh	24	821- 00- 0265	Scientist	19114.12	
3	0x160d	CUS_0xd40	December	Aaron Maashoh	24	821- 00- 0265	Scientist	19114.12	
4	0x1616	CUS_0x21b1	September	Rick Rothackerj	28	004- 07- 5839	_____	34847.84	

5 rows × 27 columns



In [157]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50000 entries, 0 to 49999
Data columns (total 27 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   ID                                    50000 non-null  object
1   Customer_ID                          50000 non-null  object
2   Month                                50000 non-null  object
3   Name                                  44985 non-null  object
4   Age                                   50000 non-null  object
5   SSN                                   50000 non-null  object
6   Occupation                           50000 non-null  object
7   Annual_Income                        50000 non-null  object
8   Monthly_Inhand_Salary                42502 non-null  float64
9   Num_Bank_Accounts                    50000 non-null  int64
10  Num_Credit_Card                       50000 non-null  int64
11  Interest_Rate                        50000 non-null  int64
12  Num_of_Loan                           50000 non-null  object
13  Type_of_Loan                          44296 non-null  object
14  Delay_from_due_date                  50000 non-null  int64
15  Num_of_Delayed_Payment               46502 non-null  object
16  Changed_Credit_Limit                 50000 non-null  object
17  Num_Credit_Inquiries                 48965 non-null  float64
18  Credit_Mix                           50000 non-null  object
19  Outstanding_Debt                     50000 non-null  object
20  Credit_Utilization_Ratio             50000 non-null  float64
21  Credit_History_Age                   45530 non-null  object
22  Payment_of_Min_Amount                 50000 non-null  object
23  Total_EMI_per_month                  50000 non-null  float64
24  Amount_invested_monthly              47729 non-null  object
25  Payment_Behaviour                     50000 non-null  object
26  Monthly_Balance                       49438 non-null  object
dtypes: float64(4), int64(4), object(19)
memory usage: 10.3+ MB
```

In [158]: df.describe()

Out[158]:

	Monthly_Inhand_Salary	Num_Bank_Accounts	Num_Credit_Card	Interest_Rate	Delay_from
count	42502.000000	50000.000000	50000.000000	50000.000000	50000.000000
mean	4182.004291	16.838260	22.921480	68.772640	10.000000
std	3174.109304	116.396848	129.314804	451.602363	10.000000
min	303.645417	-1.000000	0.000000	1.000000	0.000000
25%	1625.188333	3.000000	4.000000	8.000000	0.000000
50%	3086.305000	6.000000	5.000000	13.000000	0.000000
75%	5934.189094	7.000000	7.000000	20.000000	0.000000
max	15204.633333	1798.000000	1499.000000	5799.000000	10.000000

remove outliers

```
In [159]: df.isnull().sum()
```

```
Out[159]: ID                                0
Customer_ID                               0
Month                                    0
Name                                    5015
Age                                      0
SSN                                      0
Occupation                               0
Annual_Income                           0
Monthly_Inhand_Salary                   7498
Num_Bank_Accounts                       0
Num_Credit_Card                         0
Interest_Rate                           0
Num_of_Loan                             0
Type_of_Loan                           5704
Delay_from_due_date                     0
Num_of_Delayed_Payment                   3498
Changed_Credit_Limit                    0
Num_Credit_Inquiries                    1035
Credit_Mix                             0
Outstanding_Debt                        0
Credit_Utilization_Ratio                0
Credit_History_Age                     4470
Payment_of_Min_Amount                   0
Total_EMI_per_month                     0
Amount_invested_monthly                 2271
Payment_Behaviour                       0
Monthly_Balance                         562
dtype: int64
```

```
In [160]: df.duplicated().sum()
```

```
Out[160]: 0
```

```
In [161]: df.columns
```

```
Out[161]: Index(['ID', 'Customer_ID', 'Month', 'Name', 'Age', 'SSN', 'Occupation',
                  'Annual_Income', 'Monthly_Inhand_Salary', 'Num_Bank_Accounts',
                  'Num_Credit_Card', 'Interest_Rate', 'Num_of_Loan', 'Type_of_Loan',
                  'Delay_from_due_date', 'Num_of_Delayed_Payment', 'Changed_Credit_Limit',
                  'Num_Credit_Inquiries', 'Credit_Mix', 'Outstanding_Debt',
                  'Credit_Utilization_Ratio', 'Credit_History_Age',
                  'Payment_of_Min_Amount', 'Total_EMI_per_month',
                  'Amount_invested_monthly', 'Payment_Behaviour', 'Monthly_Balance'],
                  dtype='object')
```

```
In [162]: df.Occupation.value_counts()
```

```
Out[162]: Occupation
_____
Lawyer      3438
Engineer    3324
Architect   3212
Mechanic     3195
Developer    3168
Accountant   3146
Media_Manager 3133
Scientist    3130
Teacher      3103
Entrepreneur 3103
Journalist   3037
Doctor       3027
Manager      3000
Musician     2947
Writer       2933
Name: count, dtype: int64
```

```
In [163]: df.Num_Bank_Accounts.value_counts()
```

```
Out[163]: Num_Bank_Accounts
6         6504
7         6408
8         6387
4         6100
5         6068
...
1247      1
1721      1
703       1
1500      1
640       1
Name: count, Length: 540, dtype: int64
```

```
In [164]: df.Credit_Utilization_Ratio.value_counts()
```

```
Out[164]: Credit_Utilization_Ratio
34.108530    1
35.030402    1
33.053114    1
33.811894    1
32.430559    1
..
29.566123    1
38.135424    1
38.226475    1
31.291849    1
29.150995    1
Name: count, Length: 50000, dtype: int64
```

```
In [165]: df.Credit_History_Age.value_counts()
```

```
Out[165]: Credit_History_Age
16 Years and 1 Months    254
20 Years and 1 Months    254
18 Years and 7 Months    252
19 Years and 7 Months    252
18 Years and 6 Months    250
...
4 Years and 5 Months     21
0 Years and 11 Months    16
33 Years and 11 Months   15
34 Years and 0 Months    14
0 Years and 10 Months    13
Name: count, Length: 399, dtype: int64
```

Data Exploration and Preprocessing:

```
In [166]: df.rename(columns={'Credit_Mix':'Credit Scores'}, inplace=True)
```

```
In [167]: df.head()
```

```
Out[167]:
```

	ID	Customer_ID	Month	Name	Age	SSN	Occupation	Annual_Income	Monthly
0	0x160a	CUS_0xd40	September	Aaron Maashoh	23	821-00-0265	Scientist	19114.12	
1	0x160b	CUS_0xd40	October	Aaron Maashoh	24	821-00-0265	Scientist	19114.12	
2	0x160c	CUS_0xd40	November	Aaron Maashoh	24	821-00-0265	Scientist	19114.12	
3	0x160d	CUS_0xd40	December	Aaron Maashoh	24	821-00-0265	Scientist	19114.12	
4	0x1616	CUS_0x21b1	September	Rick Rothackerj	28	004-07-5839		34847.84	

5 rows × 27 columns

```
In [168]: df = df.drop(['ID', 'Age', 'Customer_ID', 'Name', 'SSN', 'Month', 'Monthly_Inhand_Sa
```

In [169]: `df.head()`

Out[169]:

	Occupation	Annual_Income	Num_Bank_Accounts	Num_Credit_Card	Interest_Rate	Num_of_L
0	Scientist	19114.12	3	4	3	
1	Scientist	19114.12	3	4	3	
2	Scientist	19114.12	3	4	3	
3	Scientist	19114.12	3	4	3	
4	_____	34847.84	2	4	6	

Handle any missing values

In [170]: *#Credit Scores this is our target variable*
`df['Credit Scores'].unique()`

Out[170]: `array(['Good', '_', 'Standard', 'Bad'], dtype=object)`

In [171]:
`df['Credit Scores'] = df['Credit Scores'].replace('_', pd.NA)`

In [172]: `df.isnull().sum()`

Out[172]:

Occupation	0
Annual_Income	0
Num_Bank_Accounts	0
Num_Credit_Card	0
Interest_Rate	0
Num_of_Loan	0
Delay_from_due_date	0
Changed_Credit_Limit	0
Num_Credit_Inquiries	1035
Credit Scores	9805
Outstanding_Debt	0
Credit_Utilization_Ratio	0
Payment_of_Min_Amount	0
Total_EMI_per_month	0
Payment_Behaviour	0
Monthly_Balance	562
dtype:	int64

```
In [173]: df.isnull().sum()
```

```
Out[173]: Occupation                0
Annual_Income                    0
Num_Bank_Accounts                0
Num_Credit_Card                  0
Interest_Rate                    0
Num_of_Loan                      0
Delay_from_due_date              0
Changed_Credit_Limit             0
Num_Credit_Inquiries             1035
Credit_Scores                    9805
Outstanding_Debt                 0
Credit_Utilization_Ratio         0
Payment_of_Min_Amount            0
Total_EMI_per_month              0
Payment_Behaviour                0
Monthly_Balance                  562
dtype: int64
```

```
In [174]: df.isnull().sum()
```

```
Out[174]: Occupation                0
Annual_Income                    0
Num_Bank_Accounts                0
Num_Credit_Card                  0
Interest_Rate                    0
Num_of_Loan                      0
Delay_from_due_date              0
Changed_Credit_Limit             0
Num_Credit_Inquiries             1035
Credit_Scores                    9805
Outstanding_Debt                 0
Credit_Utilization_Ratio         0
Payment_of_Min_Amount            0
Total_EMI_per_month              0
Payment_Behaviour                0
Monthly_Balance                  562
dtype: int64
```

```
In [175]: df.dropna(inplace=True)
```

```
In [176]: df[['Credit_Scores']].isnull().sum()
```

```
Out[176]: Credit_Scores    0
dtype: int64
```

```
In [177]: df.isnull().sum()
```

```
Out[177]: Occupation                0
Annual_Income                      0
Num_Bank_Accounts                  0
Num_Credit_Card                    0
Interest_Rate                      0
Num_of_Loan                        0
Delay_from_due_date                0
Changed_Credit_Limit               0
Num_Credit_Inquiries               0
Credit_Scores                     0
Outstanding_Debt                   0
Credit_Utilization_Ratio           0
Payment_of_Min_Amount              0
Total_EMI_per_month                0
Payment_Behaviour                  0
Monthly_Balance                    0
dtype: int64
```

```
In [178]: df.shape
```

```
Out[178]: (38912, 16)
```

```
In [179]: sns.set()
```

Type *Markdown* and LaTeX: α^2

outliers finding


```
In [180]: for col in df.columns:
            if is_numeric_dtype(df[col]):
                sns.boxplot(data=df[col])
                plt.title(f'Box plot for {col}')
                plt.show()
```

Box plot for Num_Bank_Accounts



```
In [181]: df.describe()
```


```
Out[181]:
```

	Num_Bank_Accounts	Num_Credit_Card	Interest_Rate	Delay_from_due_date	Num_Credit_
count	38912.000000	38912.000000	38912.000000	38912.000000	3891
mean	16.956389	22.653963	70.375514	20.988230	2
std	117.373909	128.218919	457.446740	14.810713	19
min	-1.000000	0.000000	1.000000	-5.000000	
25%	3.000000	4.000000	7.000000	10.000000	
50%	6.000000	5.000000	13.000000	18.000000	
75%	7.000000	7.000000	20.000000	28.000000	1
max	1798.000000	1499.000000	5799.000000	67.000000	259

****if our performance is too bad then we will remove ****

```
In [182]: sns.countplot(x = df.Num_Bank_Accounts)
```

```
Out[182]: <Axes: xlabel='Num_Bank_Accounts', ylabel='count'>
```



- Encode

```
In [183]: df.columns
```

```
Out[183]: Index(['Occupation', 'Annual_Income', 'Num_Bank_Accounts', 'Num_Credit_Card',  
                'Interest_Rate', 'Num_of_Loan', 'Delay_from_due_date',  
                'Changed_Credit_Limit', 'Num_Credit_Inquiries', 'Credit_Scores',  
                'Outstanding_Debt', 'Credit_Utilization_Ratio', 'Payment_of_Min_Amount',  
                'Total_EMI_per_month', 'Payment_Behaviour', 'Monthly_Balance'],  
              dtype='object')
```

```
In [184]: df['Credit_Scores'].unique()
```

```
Out[184]: array(['Good', 'Standard', 'Bad'], dtype=object)
```

```
In [185]: from sklearn.preprocessing import LabelEncoder
for i in df.columns:
    if is_numeric_dtype(df[i]):
        continue
    else:
        df[i] = LabelEncoder().fit_transform(df[i])
```

```
In [186]: df.head()
```

```
Out[186]:
```

	Occupation	Annual_Income	Num_Bank_Accounts	Num_Credit_Card	Interest_Rate	Num_of_L
0	12	4806	3	4	3	
1	12	4806	3	4	3	
2	12	4806	3	4	3	
3	12	4806	3	4	3	
4	15	8269	2	4	6	

features selection

```
In [187]: X = df.drop('Credit Scores' , axis =1 )
Y = df[['Credit Scores']]
```

```
In [188]: xtrain , xtest , ytrain , ytest = train_test_split(X , Y , test_size=.2 , rand
```

```
In [190]: from sklearn.ensemble import ExtraTreesClassifier
etc = ExtraTreesClassifier(random_state=100)
etc.fit(xtrain , ytrain)
```

```
Out[190]: ExtraTreesClassifier(random_state=100)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [191]: etc.feature_importances_
```

```
Out[191]: array([0.0258809 , 0.03549286, 0.05967977, 0.03193272, 0.05873011,
0.12569818, 0.17731494, 0.07074362, 0.0320701 , 0.06985645,
0.02224756, 0.20690081, 0.02394781, 0.01877161, 0.04073258])
```

```
In [192]: imp = pd.DataFrame(etc.feature_importances_ , columns=['Gain_Score'])
cols = pd.DataFrame(df.columns , columns=['Feature_Names'])
gains = pd.concat([imp , cols] , axis = 1)
```

```
In [193]: gains.nlargest(12, 'Gain_Score')
```

```
Out[193]:
```

	Gain_Score	Feature_Names
11	0.206901	Credit_Utilization_Ratio
6	0.177315	Delay_from_due_date
5	0.125698	Num_of_Loan
7	0.070744	Changed_Credit_Limit
9	0.069856	Credit Scores
2	0.059680	Num_Bank_Accounts
4	0.058730	Interest_Rate
14	0.040733	Payment_Behaviour
1	0.035493	Annual_Income
8	0.032070	Num_Credit_Inquiries
3	0.031933	Num_Credit_Card
0	0.025881	Occupation

Occupation', 'Annual_Income', 'Num_Bank_Accounts', 'Num_Credit_Card', 'Interest_Rate', 'Num_of_Loan', 'Delay_from_due_date', 'Changed_Credit_Limit', 'Num_Credit_Inquiries', 'Credit Scores', 'Outstanding_Debt', 'Credit_Utilization_Ratio', 'Payment_of_Min_Amount', 'Total_EMI_per_month', 'Payment_Behaviour', 'Monthly_Balance'], dtype='object'

```
In [194]: df.shape
```

```
Out[194]: (38912, 16)
```

```
In [195]: df.columns
```

```
Out[195]: Index(['Occupation', 'Annual_Income', 'Num_Bank_Accounts', 'Num_Credit_Card',
                'Interest_Rate', 'Num_of_Loan', 'Delay_from_due_date',
                'Changed_Credit_Limit', 'Num_Credit_Inquiries', 'Credit Scores',
                'Outstanding_Debt', 'Credit_Utilization_Ratio', 'Payment_of_Min_Amount',
                'Total_EMI_per_month', 'Payment_Behaviour', 'Monthly_Balance'],
                dtype='object')
```

```
In [196]: df.drop('Occupation' , axis =1 , inplace = True)
```

Type *Markdown* and LaTeX: α^2

In [197]: `df.describe()`

Out[197]:

	Annual_Income	Num_Bank_Accounts	Num_Credit_Card	Interest_Rate	Num_of_Loan	Del
count	38912.000000	38912.000000	38912.000000	38912.000000	38912.000000	
mean	7661.399337	16.956389	22.653963	70.375514	97.001619	
std	4398.531779	117.373909	128.218919	457.446740	63.721398	
min	0.000000	-1.000000	0.000000	1.000000	0.000000	
25%	3861.750000	3.000000	4.000000	7.000000	3.000000	
50%	7692.000000	6.000000	5.000000	13.000000	102.000000	
75%	11444.000000	7.000000	7.000000	20.000000	141.000000	
max	15298.000000	1798.000000	1499.000000	5799.000000	219.000000	

In [198]: `from sklearn.preprocessing import MaxAbsScaler`

In [199]: `scaler = MaxAbsScaler()`

In [200]: `df['Annual_Income'],df['Changed_Credit_Limit'] , df['Num_Credit_Inquiries'] ,`

In [201]: `df.head()`

Out[201]:

	Annual_Income	Num_Bank_Accounts	Num_Credit_Card	Interest_Rate	Num_of_Loan	Delay_f
0	0.417949	3	4	3	121	
1	0.417949	3	4	3	121	
2	0.417949	3	4	3	121	
3	0.417949	3	4	3	121	
4	0.802816	2	4	6	3	

Model Selection:

In [202]: `from sklearn.ensemble import RandomForestClassifier , ExtraTreesClassifier ,Ad
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
import warnings as w
w.filterwarnings('ignore')`

```
In [203]: df.columns
```

```
Out[203]: Index(['Annual_Income', 'Num_Bank_Accounts', 'Num_Credit_Card',  
                'Interest_Rate', 'Num_of_Loan', 'Delay_from_due_date',  
                'Changed_Credit_Limit', 'Num_Credit_Inquiries', 'Credit_Scores',  
                'Outstanding_Debt', 'Credit_Utilization_Ratio', 'Payment_of_Min_Amount',  
                'Total_EMI_per_month', 'Payment_Behaviour', 'Monthly_Balance'],  
              dtype='object')
```

```
In [204]: X = df.drop('Credit_Scores' , axis =1 )  
Y = df[['Credit_Scores']]
```

```
In [205]: xtrain , xtest , ytrain , ytest = train_test_split(X , Y , test_size=.2 , rand
```

```
In [206]: #AdaBoostClassifier  
ada = AdaBoostClassifier()  
ada.fit(xtrain , ytrain )
```

```
Out[206]: AdaBoostClassifier()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [207]: ada.score(xtrain , ytrain)
```

```
Out[207]: 0.817790484756979
```

```
In [208]: ada.score(xtest , ytest)
```

```
Out[208]: 0.8171656173711936
```

```
In [209]: #RandomForestClassifier  
ran = RandomForestClassifier(random_state=100)
```

```
In [210]: ran.fit(xtrain , ytrain )
```

```
Out[210]: RandomForestClassifier(random_state=100)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [211]: ran.score(xtest , ytest)
```

```
Out[211]: 0.9509186688937428
```

```
In [212]: ran.score(xtrain , ytrain)
```

```
Out[212]: 1.0
```

```
In [213]: #ExtraTreesClassifier  
etc = ExtraTreesClassifier(random_state=100)
```

```
In [214]: etc.fit(xtrain , ytrain)
```

```
Out[214]: ExtraTreesClassifier(random_state=100)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [215]: etc.score(xtrain , ytrain)
```

```
Out[215]: 1.0
```

```
In [216]: etc.score(xtest , ytest)
```

```
Out[216]: 0.9225234485416934
```

```
In [217]: reg = LogisticRegression(random_state=100)
```

```
In [218]: reg.fit(xtrain , ytrain)
```

```
Out[218]: LogisticRegression(random_state=100)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [219]: reg.score(xtest , ytest)
```

```
Out[219]: 0.5255043042528588
```

```
In [220]: reg.score(xtrain , ytrain)
```

```
Out[220]: 0.5186803302386842
```

Model Training

```
In [221]: from sklearn.metrics import accuracy_score , classification_report , confusion_
```

```
In [222]: print(classification_report(ytest , etc.predict(xtest)))
```

	precision	recall	f1-score	support
0	0.91	0.94	0.92	1826
1	0.91	0.96	0.94	2367
2	0.94	0.89	0.91	3590
accuracy			0.92	7783
macro avg	0.92	0.93	0.92	7783
weighted avg	0.92	0.92	0.92	7783

```
In [223]: print(classification_report(ytest , ran.predict(xtest)))
```

	precision	recall	f1-score	support
0	0.92	0.98	0.95	1826
1	0.96	0.97	0.96	2367
2	0.97	0.93	0.95	3590
accuracy			0.95	7783
macro avg	0.95	0.96	0.95	7783
weighted avg	0.95	0.95	0.95	7783

```
In [224]: confusion_matrix(ytest , ran.predict(xtest))
```

```
Out[224]: array([[1787,    0,   39],
                 [    0, 2289,   78],
                 [ 161,  104, 3325]])
```

```
In [225]: confusion_matrix(ytest , etc.predict(xtest))
```

```
Out[225]: array([[1717,    0,  109],
                 [    0, 2266,  101],
                 [ 180,  213, 3197]])
```

Hyperparameter Tuning

```
In [ ]: #Define hyperparameters for Random Forest
rf_params = {
    'n_estimators': [101, 151, 201, 251, 301], #https://numpy.org/doc/stable/r
    'criterion': ['gini', 'entropy'],
    'max_depth': [None, 2,4,6,8,10,],
    'min_samples_split': [2,3,4,5,6,7,8,10],
    'min_samples_leaf': [1, 2,3, 4],
    'max_features': ['auto', 'sqrt', 'log2']
}
```



```
In [ ]: from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
```

```
In [ ]: rf_random_search = RandomizedSearchCV(estimator=RandomForestClassifier(random_  
                                              param_distributions=rf_params, n_iter=10  
rf_random_search.fit(xtrain, ytrain)
```

Type *Markdown* and LaTeX: α^2

```
In [ ]: rf_random_search.best_params_
```

```
In [ ]: #laptop problem korsa Hyperparameter Tuning r kora galo na
```